STAT 331 Final Project

Krishna Prem Pasumarthy & Islam Amin

April 12, 2020

1   Summary

2   Descriptive Statistics

First, take a look at summary statistics of the `fhsd` dataset.

Table 1: Summary Statistics

| chdrisk	sex	totchol	age	sysbp	diabp	cursmoke	cigpday	bmi	diabetes	bpmeds	heartrte	glucose	prevmi	prevstrk	prevhyp	hdlc	ldlc
Min. :0.0050	Female:1305	Min. :112.0	Min. :44.00	Min. : 86.0	Min. : 30.00	No :1504	Min. : 0.00	Min. :14.43	No :2142	No :1973	Min. : 44.00	Min. : 46.00	No :2189	No :2260	No : 957	Min. : 10.00	Min. : 20.0
1st Qu.:0.1320	Male :1001	1st Qu.:207.0	1st Qu.:53.00	1st Qu.:122.5	1st Qu.: 73.00	Yes: 802	1st Qu.: 0.00	1st Qu.:23.22	Yes: 164	Yes: 333	1st Qu.: 70.00	1st Qu.: 75.00	Yes: 117	Yes: 46	Yes:1349	1st Qu.: 38.00	1st Qu.:152.0
Median :0.2240	NA	Median :235.5	Median :60.00	Median :136.0	Median : 80.00	NA	Median : 0.00	Median :25.40	NA	NA	Median : 76.00	Median : 83.00	NA	NA	NA	Median : 47.00	Median :180.0
Mean :0.2655	NA	Mean :237.8	Mean :60.23	Mean :139.2	Mean : 81.07	NA	Mean : 6.84	Mean :25.78	NA	NA	Mean : 77.61	Mean : 89.07	NA	NA	NA	Mean : 48.89	Mean :183.1
3rd Qu.:0.3448	NA	3rd Qu.:265.0	3rd Qu.:67.00	3rd Qu.:153.0	3rd Qu.: 88.00	NA	3rd Qu.:10.00	3rd Qu.:27.91	NA	NA	3rd Qu.: 85.00	3rd Qu.: 95.00	NA	NA	NA	3rd Qu.: 57.00	3rd Qu.:210.0
Max. :0.9770	NA	Max. :625.0	Max. :81.00	Max. :246.0	Max. :130.00	NA	Max. :80.00	Max. :46.52	NA	NA	Max. :150.00	Max. :478.00	NA	NA	NA	Max. :189.00	Max. :565.0

Then take a look at `chdrisk` grouped by `sex` as well as `chdrisk` grouped by `cursmoke`.

```
## fhsd$sex: Female
##    Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##   0.005   0.104   0.179   0.215   0.285   0.949
## -----------------------------------------------------------
## fhsd$sex: Male
##    Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##  0.0210  0.1860  0.2860  0.3314  0.4060  0.9770

## fhsd$cursmoke: No
##    Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##  0.0050  0.1390  0.2350  0.2754  0.3580  0.9770
## -----------------------------------------------------------
## fhsd$cursmoke: Yes
##    Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##  0.0080  0.1220  0.1995  0.2471  0.3140  0.9710

## fhsd$prevstrk: No
##    Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##  0.0050  0.1300  0.2200  0.2611  0.3392  0.9770
## -----------------------------------------------------------
## fhsd$prevstrk: Yes
##    Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##  0.2020  0.3412  0.4410  0.4820  0.5060  0.9660

## fhsd$prevhyp: No
##    Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##   0.005   0.077   0.140   0.176   0.216   0.944
## -----------------------------------------------------------
## fhsd$prevhyp: Yes
##    Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##  0.0320  0.1980  0.2890  0.3291  0.4010  0.9770

## fhsd$prevmi: No
##    Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##  0.0050  0.1270  0.2120  0.2309  0.3220  0.6920
```

```
## ------------------------------------------------------------
## fhsd$prevmi: Yes
##     Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##  0.8240  0.8940  0.9150  0.9133  0.9360  0.9770
```

[ADD SOME COMMENTS HERE REGARDING SUMMARY]

Now take a look at pair plots of all numeric explanatory variates i.e. variates excluding response variate `chdrisk` and logical variates such as `cursmoke`.

Pair Plots of Continuous Variates
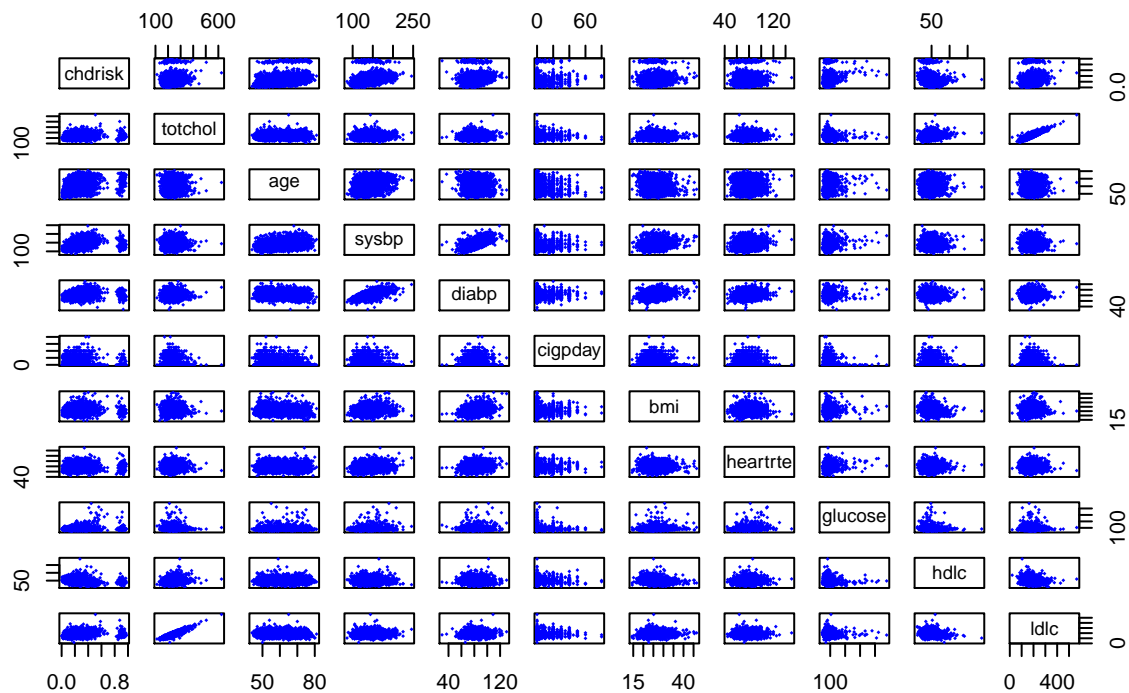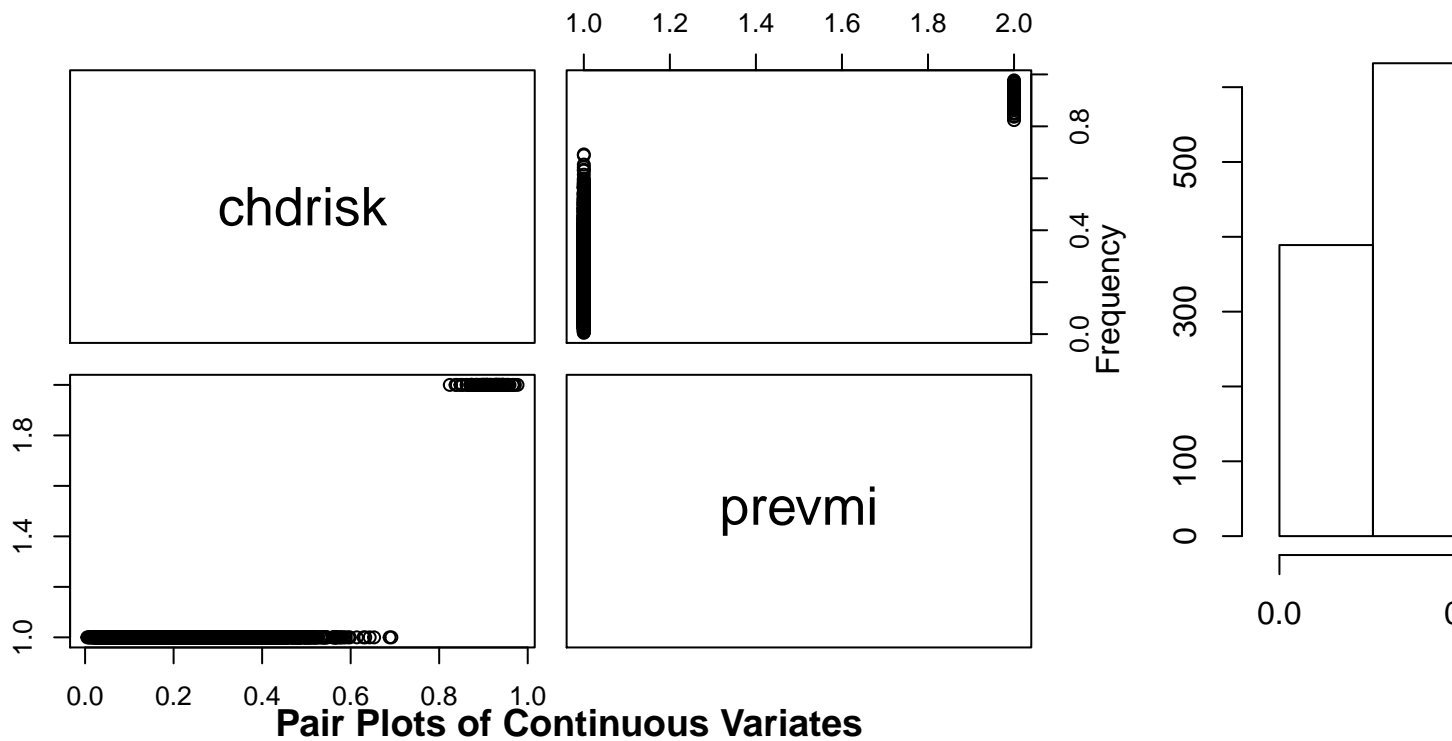
From the pair plots, we can observe a strong correlation between low density lipoprotein cholesterol and serum total cholestrol. This correlation could be explained by the fact that there could be a relationship between the amount [TO BE CONTINUED]

Now take a look at the VIFs of these variates.

```
##     sexMale     totchol         age       sysbp       diabp cursmokeYes
##    1.225191   10.634882    1.489926    2.918660    2.406411    2.978609
##     cigpday         bmi diabetesYes   bpmedsYes    heartrte     glucose
##    2.973594    1.181865    1.286401    1.214744    1.105902    1.308923
##   prevmiYes prevstrkYes  prevhypYes        hdlc        ldlc
##    1.067134    1.045746    1.823014    2.287571   10.367649
```

[ADD COMMENTS]

3 Candidate Models

3.1 Automated Model Selection

```r
library(gtools)

load_calcs = TRUE
# model with only intercept
M0 <- lm(I(logit(chdrisk)) ~ 1, data = fhsd)
Mmax <- lm(I(logit(chdrisk)) ~ (.)^2, data = fhsd)
# starting model for stepwise selection
Mstart <- lm(I(logit(chdrisk)) ~ ., data = fhsd)

# find model coefficients which are NA
beta.max <- coef(Mmax)
names(beta.max)[is.na(beta.max)]
```

```
## [1] "cursmokeYes:cigpday"  "bpmedsYes:prevhypYes"
```

```r
# find the problem with the NA coeffs
kable(table(fhsd[c("cursmoke", "cigpday")]), "latex")
```

	0	1	2	3	4	5	6	7	8	9	10	12	14	15	16	17	18	19	20	23	25	26	27
No	1504	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
Yes	0	16	18	34	11	18	24	9	18	5	76	3	3	50	6	1	8	1	279	1	14	1	1

```r
kable(table(fhsd[c("bpmeds", "prevhyp")]), "latex")
```

	No	Yes
No	957	1016
Yes	0	333

```r
# remove the coeffs with the problem and add quadratic terms for the continuous variables
Mmax <- lm(I(logit(chdrisk)) ~ (.)^2 - cursmoke:cigpday - bpmeds:prevhyp +
             I(totchol ^ 2) + I(sysbp ^ 2) + I(diabp ^ 2)
           + I(bmi ^ 2) + I(glucose ^ 2)
           + I(hdlc ^ 2) + I(ldlc ^ 2), data = fhsd)

anyNA(coef(Mmax)) # check if there are any remaining NAs
```

```
## [1] FALSE
```

```r
if(!load_calcs){

  #forward model selection
  system.time({
    Mfwd <- step(object = M0,
                 scope = list(lower = M0, upper = Mmax),
                 direction = "forward", trace = FALSE)
  })

  #backward model selection
  system.time({
    Mback <- step(object = Mmax,
                  scope = list(lower = M0, upper = Mmax),
                  direction = "backward", trace = FALSE)
  })
```

```r
  #stepwise model selection
  system.time({
    Mstep <- step(object = Mstart,
                  scope = list(lower = M0, upper = Mmax),
                  direction = "both", trace = FALSE)
  })

}


# the caching/loading block
if(!load_calcs) {
  saveRDS(list(Mfwd = Mfwd, Mback = Mback, Mstep = Mstep), file = "models_automated.rds")
} else {
  # just load the calculations
  tmp <- readRDS("models_automated.rds")
  Mfwd <- tmp$Mfwd
  Mback <- tmp$Mback
  Mstep <- tmp$Mstep
  rm(tmp) # optionally remove tmp from workspace
}


# Stepwise model selection
Mstep$call
```

```
## lm(formula = I(logit(chdrisk)) ~ sex + totchol + age + sysbp +
##     diabp + cursmoke + cigpday + bmi + diabetes + bpmeds + heartrte +
##     glucose + prevmi + prevstrk + prevhyp + hdlc + ldlc + I(hdlc^2) +
##     I(bmi^2) + I(diabp^2) + I(sysbp^2) + sysbp:prevmi + totchol:prevhyp +
##     diabetes:prevmi + prevhyp:ldlc + sysbp:prevhyp + totchol:heartrte +
##     sysbp:diabetes + diabp:bmi + diabp:hdlc + prevmi:hdlc + prevmi:prevhyp +
##     sex:glucose + age:ldlc + age:heartrte + cigpday:hdlc + bmi:ldlc +
##     totchol:hdlc + totchol:prevmi + sysbp:heartrte + sysbp:bpmeds +
##     cursmoke:hdlc + prevmi:prevstrk + diabetes:hdlc + sex:sysbp +
##     cigpday:glucose + heartrte:glucose + diabp:glucose + cursmoke:ldlc +
##     age:cigpday + age:hdlc + hdlc:ldlc + age:prevhyp + diabp:prevhyp +
##     diabp:cursmoke + diabp:cigpday + bmi:bpmeds + bpmeds:glucose +
##     age:prevmi + sex:ldlc + cigpday:heartrte + cigpday:prevmi +
##     glucose:prevmi + heartrte:prevmi + bpmeds:prevstrk, data = fhsd)
```

```r
# Forward model selection
Mfwd$call
```

```
## lm(formula = I(logit(chdrisk)) ~ prevmi + sysbp + sex + age +
##     ldlc + prevhyp + diabetes + hdlc + I(hdlc^2) + cigpday +
##     I(bmi^2) + bmi + totchol + I(glucose^2) + I(sysbp^2) + bpmeds +
##     heartrte + cursmoke + prevstrk + prevmi:sysbp + sysbp:age +
##     prevhyp:hdlc + prevmi:diabetes + sysbp:prevhyp + prevhyp:totchol +
##     sysbp:diabetes + prevmi:hdlc + prevmi:prevhyp + age:ldlc +
##     age:cigpday + hdlc:cigpday + prevhyp:bmi + ldlc:bmi + prevmi:totchol +
##     ldlc:prevhyp + sysbp:bpmeds + sysbp:hdlc + hdlc:totchol +
##     totchol:heartrte + age:heartrte + diabetes:hdlc + sysbp:heartrte +
##     bmi:bpmeds + sysbp:sex + ldlc:hdlc + prevmi:bmi + age:bmi +
```

```
##       prevmi:age + sysbp:cursmoke + hdlc:cursmoke + ldlc:cursmoke +
##       prevmi:cigpday + sex:diabetes + prevmi:prevstrk, data = fhsd)
```

```r
# Backward model selection
Mback$call
```

```
## lm(formula = I(logit(chdrisk)) ~ sex + totchol + age + sysbp +
##       diabp + cursmoke + cigpday + bmi + diabetes + bpmeds + heartrte +
##       glucose + prevmi + prevstrk + prevhyp + hdlc + ldlc + I(totchol^2) +
##       I(sysbp^2) + I(diabp^2) + I(bmi^2) + I(hdlc^2) + I(ldlc^2) +
##       sex:totchol + sex:sysbp + sex:glucose + sex:prevstrk + sex:prevhyp +
##       totchol:age + totchol:bpmeds + totchol:heartrte + totchol:prevmi +
##       totchol:prevstrk + totchol:prevhyp + totchol:hdlc + totchol:ldlc +
##       age:cursmoke + age:bmi + age:heartrte + age:prevmi + age:prevhyp +
##       age:hdlc + sysbp:diabetes + sysbp:bpmeds + sysbp:heartrte +
##       sysbp:prevmi + sysbp:prevhyp + diabp:cursmoke + diabp:cigpday +
##       diabp:bmi + diabp:glucose + diabp:prevhyp + diabp:hdlc +
##       cursmoke:bmi + cursmoke:hdlc + cursmoke:ldlc + cigpday:bmi +
##       cigpday:heartrte + cigpday:glucose + cigpday:prevmi + cigpday:hdlc +
##       bmi:prevmi + bmi:prevhyp + bmi:ldlc + diabetes:prevmi + diabetes:hdlc +
##       bpmeds:glucose + bpmeds:prevstrk + bpmeds:ldlc + heartrte:glucose +
##       heartrte:prevmi + glucose:prevmi + prevmi:prevhyp + prevmi:hdlc +
##       prevhyp:ldlc, data = fhsd)
```

```r
beta.fwd = coef(Mfwd)
beta.back = coef(Mback)
beta.step = coef(Mstep)

identical(names(beta.fwd)[names(beta.fwd) %in% names(beta.back)], names(beta.fwd))
```

```
## [1] FALSE
```

```r
identical(names(beta.fwd)[names(beta.fwd) %in% names(beta.step)], names(beta.fwd))
```

```
## [1] FALSE
```

```r
identical(names(beta.back)[names(beta.back) %in% names(beta.step)], names(beta.back))
```

```
## [1] FALSE
```

3.2   Manual Model Selection

```r
library(stringr) # For string operations
```

```
## Warning: package 'stringr' was built under R version 3.5.2
```

```r
table <- c() # Initialize empty vector

names.table <- names(beta.step)                    # Obtain variate names in stepwise model
names.table <- str_remove_all(names.table,"Yes") # Remove "Yes" from interactions
names.table <- str_remove_all(names.table,"Male") # Remove "Male"

# Perform F-tests by removing one variate at a time
  for(i in names.table){
    # Obtain model without variate i
    mdl <- lm(as.formula(paste0("update(Mstep, . ~ . -", i,")")),data = fhsd)
    test <- anova(Mstep,mdl)                # F-Test between Stepwise and reduced model
```

```r
  table <- cbind(table,test$`Pr(>F)`[2]) # Add corresponding p-value to the table
 }
table <- as.data.frame(table)
colnames(table) <- names.table # Add appropriate column names to the table

sort(table,decreasing = TRUE)   # Look at the variates that show least significance
```

```
##    cigpday:heartrte bpmeds:prevstrk bpmeds:glucose diabp:cigpday    cigpday
## 1         0.1506282       0.1492283      0.1189197     0.1155989 0.1151079
##      sex:ldlc age:prevmi cigpday:prevmi hdlc:ldlc bmi:bpmeds prevmi:prevstrk
## 1 0.1141483 0.1097987       0.1051865 0.0923568 0.0855445      0.06997763
##    heartrte:prevmi glucose:prevmi I(sysbp^2) cursmoke:hdlc age:heartrte
## 1       0.06451949     0.05883116 0.0585469    0.05660935    0.05562064
##    age:hdlc cursmoke:ldlc  sex:sysbp sysbp:bpmeds    age:ldlc
## 1 0.0510796     0.0417893 0.03623249    0.0300776 0.02915113
##    cigpday:glucose prevmi:prevhyp       hdlc sex:glucose diabetes:hdlc
## 1        0.0291137     0.02242217 0.01880445  0.01702301    0.01394662
##    diabp:glucose    bmi:ldlc totchol:hdlc    bpmeds age:cigpday
## 1     0.01362058 0.009985489  0.009840662 0.0077735 0.006735591
##    heartrte:glucose    cursmoke totchol:prevmi sysbp:heartrte diabp:prevhyp
## 1       0.004772297 0.004188557    0.003609581    0.002926201   0.001409115
##    diabp:cursmoke prevhyp:ldlc        bmi age:prevhyp sysbp:diabetes
## 1     0.001393474  0.00066789 0.0006664543 0.0005753017   0.0004931994
##    I(hdlc^2)  diabp:hdlc sysbp:prevhyp cigpday:hdlc  prevmi:hdlc
## 1 0.000320732 0.0001422969  0.0001292531 0.0001038006 7.056001e-05
##    diabetes:prevmi       diabp totchol:heartrte    diabp:bmi sysbp:prevmi
## 1     6.226049e-05 6.021714e-05     3.512093e-05 2.940165e-05 2.305381e-05
##           sex     heartrte          age totchol:prevhyp      I(bmi^2)
## 1 2.396724e-06 9.478088e-07 4.238229e-07    1.203731e-09 2.735937e-11
##     I(diabp^2)      prevmi      prevhyp
## 1 1.257752e-19 1.595006e-22 1.119628e-27
```

```r
# Remove first three variates from table since starting the 4th variate, reduction becomes significant
anova(Mstep, update(Mstep,. ~ . - cigpday:heartrte - bpmeds:prevstrk - bpmeds:glucose))
```

```
## Analysis of Variance Table
##
## Model 1: I(logit(chdrisk)) ~ sex + totchol + age + sysbp + diabp + cursmoke +
##     cigpday + bmi + diabetes + bpmeds + heartrte + glucose +
##     prevmi + prevstrk + prevhyp + hdlc + ldlc + I(hdlc^2) + I(bmi^2) +
##     I(diabp^2) + I(sysbp^2) + sysbp:prevmi + totchol:prevhyp +
##     diabetes:prevmi + prevhyp:ldlc + sysbp:prevhyp + totchol:heartrte +
##     sysbp:diabetes + diabp:bmi + diabp:hdlc + prevmi:hdlc + prevmi:prevhyp +
##     sex:glucose + age:ldlc + age:heartrte + cigpday:hdlc + bmi:ldlc +
##     totchol:hdlc + totchol:prevmi + sysbp:heartrte + sysbp:bpmeds +
##     cursmoke:hdlc + prevmi:prevstrk + diabetes:hdlc + sex:sysbp +
##     cigpday:glucose + heartrte:glucose + diabp:glucose + cursmoke:ldlc +
##     age:cigpday + age:hdlc + hdlc:ldlc + age:prevhyp + diabp:prevhyp +
##     diabp:cursmoke + diabp:cigpday + bmi:bpmeds + bpmeds:glucose +
##     age:prevmi + sex:ldlc + cigpday:heartrte + cigpday:prevmi +
##     glucose:prevmi + heartrte:prevmi + bpmeds:prevstrk
## Model 2: I(logit(chdrisk)) ~ sex + totchol + age + sysbp + diabp + cursmoke +
##     cigpday + bmi + diabetes + bpmeds + heartrte + glucose +
##     prevmi + prevstrk + prevhyp + hdlc + ldlc + I(hdlc^2) + I(bmi^2) +
```

```
##     I(diabp^2) + I(sysbp^2) + sysbp:prevmi + totchol:prevhyp +
##     diabetes:prevmi + prevhyp:ldlc + sysbp:prevhyp + totchol:heartrte +
##     sysbp:diabetes + diabp:bmi + diabp:hdlc + prevmi:hdlc + prevmi:prevhyp +
##     sex:glucose + age:ldlc + age:heartrte + cigpday:hdlc + bmi:ldlc +
##     totchol:hdlc + totchol:prevmi + sysbp:heartrte + sysbp:bpmeds +
##     cursmoke:hdlc + prevmi:prevstrk + diabetes:hdlc + sex:sysbp +
##     cigpday:glucose + heartrte:glucose + diabp:glucose + cursmoke:ldlc +
##     age:cigpday + age:hdlc + hdlc:ldlc + age:prevhyp + diabp:prevhyp +
##     diabp:cursmoke + diabp:cigpday + bmi:bpmeds + age:prevmi +
##     sex:ldlc + cigpday:prevmi + glucose:prevmi + heartrte:prevmi
##   Res.Df    RSS Df Sum of Sq      F  Pr(>F)
## 1   2240 489.70
## 2   2243 491.12 -3   -1.4219 2.1681 0.08981 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

4 Model Diagnostics

5 Model Selection

6 Discussion